

EUGENE L. BUCHANAN

Staff AI Engineer — Multi-Agent Frameworks & Agentic Infrastructure

San Quintín, BC MX (Remote / PT) · +1 (909) 545 5384 · eugene@serviceofothers.org · github.com/eugeneb50



August 3, 2026

ClickUp Hiring Team

AI Platform — Staff AI Engineer, Multi-Agent Frameworks

Via: clickup.com/careers

Re: Staff AI Engineer — Multi-Agent Frameworks (#LI-REMOTE)

Dear ClickUp AI Platform Team,

I'm applying for the Staff AI Engineer position on the Multi-Agent Frameworks team. ClickUp's vision of a converged AI workspace — unifying tasks, docs, chat, calendar, and search through context-driven AI — is the exact problem space I've been working in: building platforms that let agents coordinate, evaluate, and operate at scale.

On the ZeroClaw agentic AI gateway (32.2k stars, Rust), I shipped an observer that reports agent lifecycle state — idle, working, blocked, released — via JSON-RPC over Unix domain sockets (PR #8337, 1,132 additions). That work is multi-agent orchestration: bounded I/O threading, env-driven discovery, and reliability guards for unresponsive sockets without hanging the agent loop. I also reviewed a 6,684-line security PR adding multi-user auth, deny-by-default permission profiles, and principal isolation at RPC dispatch (PR #8672) — the same isolation problem multi-agent platforms face.

At Knowledgecity I was Product Owner for Integrations, shipping SAP, Oracle, Workday, UKG, Coursera, and Zoom via SAML, OAuth, SFTP, and custom APIs on AWS/Postgres/React. Orchestrating six enterprise services with different auth models and error semantics is the same coordination challenge multi-agent frameworks solve. I also built the evaluation framework (Cypress, Selenium, JUnit, health dashboards, alerting) for testing those integrations in complex scenarios, and applied AI prompt engineering and chatbot tooling in production.

Here's how I map to your qualifications:

- **Multiple LLMs & orchestration.** Integrating OpenAI, Anthropic, and Gemini providers daily through Hermes Agent; agent lifecycle orchestration in production via Rust + JSON-RPC.
- **Evaluation frameworks.** 27 years building test automation and quality pipelines; four rounds of reviewer blockers resolved on the ZeroClaw PR through targeted evaluation of each failure mode.
- **Multi-agent frameworks.** Shipped herdr agent reporting integration (multi-agent coordination without requiring a herdr adapter); coordinating on a standardized generic agent integration adapter.
- **AI privacy & compliance.** Reviewed deny-by-default permission profiles, peercred, SSH challenge-response, and OIDC auth flows plus 27 years in regulated, compliance-driven environments.
- **Search technologies.** Elasticsearch, PostgreSQL full-text, and integration search across enterprise platforms at Knowledgecity and in the ZeroClaw ecosystem.

I'm AI-native in the way ClickUp means it: I use AI agents daily, I build the infrastructure they run on, and I evaluate their behavior under real-world complexity. I'd welcome the chance to bring that to ClickUp's AI Platform team and help build the future of converged work. Happy to walk through any of the PRs or code in a technical interview.

Best regards,

Eugene L. Buchanan

Staff AI Engineer — Multi-Agent Frameworks

eugene@serviceofothers.org · +1 (909) 545 5384 · github.com/eugeneb50