

EUGENE L. BUCHANAN

Staff AI Engineer — Multi-Agent Frameworks & Agentic Infrastructure

Apple Valley, CA (Remote) · +1 (909) 545 5384 · eugene@serviceofothers.org · github.com/eugeneb50



August 11, 2026

ClickUp Hiring Team

AI Platform — Staff AI Engineer, Multi-Agent Frameworks

Via: clickup.com/careers

Re: Staff AI Engineer — Multi-Agent Frameworks (#LI-REMOTE)

Dear ClickUp AI Platform Team,

I'm applying for the Staff AI Engineer — Multi-Agent Frameworks role on your AI Platform team. I build backend platforms where users create, deploy, and coordinate intelligent agents. That work needs full context for humans and agents to act side by side — exactly the problem I've been solving in production.


On ZeroClaw's 32.2k-star Rust gateway, I ship multi-agent orchestration. An observer reports agent lifecycle state — idle, working, blocked, released — through JSON-RPC over Unix domain sockets (PR #8337). A merged context-window meter unifies 9 LLM providers — OpenAI, Anthropic, Cohere, and more — into one source of truth. That drives model routing and cost attribution (PR #7946). I prototype workflows with LangGraph, build MCP servers in Rust, and cover the backend stack (Rust, Python, Node on PostgreSQL/AWS).

At Knowledgecity, I owned integrations as Product Owner. I shipped SAP, Oracle, Workday, UKG, Coursera, and Zoom over SAML, OAuth, SFTP, and REST APIs. Each service uses different auth and error semantics — the same coordination puzzle that multi-agent frameworks solve. I also built evaluation tooling (Cypress, Selenium, JUnit, health dashboards, alerting) to test complex deployments at the system level. I handled AI privacy (deny-by-default permissions, principal isolation, OIDC) and tied in search (Elasticsearch, PostgreSQL).

How I fit this role:

- **Multi-Agent Frameworks & orchestration.** Agent lifecycle observer, LangGraph workflows, and an MCP server in a 32k-star Rust gateway.
- **Multiple LLMs & model routing.** OpenAI, Anthropic, Cohere, Gemini — one context window and per-provider cost attribution.
- **Evaluation frameworks.** Deep system-level testing; evaluation at both agent and system-level dynamics.
- **AI privacy & search.** Deny-by-default permissions, principal isolation; Elasticsearch and Postgres full-text.

ClickUp's Value Hiring Eugene



I'm AI-native the way ClickUp means it: I use agents daily, I build the infrastructure they run on, and I evaluate their behavior under real-world complexity. I'd welcome the chance to bring that to your AI Platform team — happy to walk through any PR in an interview.

Best regards,

Eugene L. Buchanan

Staff AI Engineer — Multi-Agent Frameworks

Apple Valley, CA · eugene@serviceofothers.org · github.com/eugeneb50