



Date: August 11, 2026

To:

ClickUp Recruiting Team Mango Technologies, Inc. (ClickUp) San Diego, CA (Remote, US)

SUBJECT: Application for Staff AI Engineer - Multi-Agent Frameworks (JID: 1203026f-5c19-45e5-a0c2-b5cc0338a1e8)

Dear ClickUp Recruiting Team,


I am writing to express my enthusiastic interest in the Staff AI Engineer - Multi-Agent Frameworks position at ClickUp. Your mission to maximize human productivity fundamentally aligns with my software engineering philosophy. While other companies scramble to bundle fragmented tools or bolt on AI as an afterthought, ClickUp is doing it correct building the world's first truly converged AI workspace unifying tasks, docs, chat, and search, all supercharged by context-driven AI. Your culture of grit, speed, and breaking the status quo (especially your core value 'Normal f*cking sucks') is precisely the high-intensity environment where I excel.

As an active contributor to the ****ZeroClaw**** Rust autonomous AI agent infrastructure and ****herdr**** multi-agent multiplexer, I have spent the last several years designing, debugging, and deploying the exact systems-level patterns ClickUp is leveraging. I specialize in high-performance LLM orchestration, context-management protocols, and rigorous evaluation frameworks. I am highly confident in my ability to immediately add value to your AI Platform team.

Here are several concrete ways my systems engineering work translates to what I can do for ClickUp:

- **Dynamic Context Window Resolution (PR #8966):** I designed and implemented live LLM provider identity propagation on usage events, dynamically resolving the active context window boundaries from the serving provider at runtime. This guarantees accurate token-use observability and prevents context overflow errors during multi-provider fallback.
- **Robust High-Performance Communication (PR #8337):** I engineered a JSON-RPC communication layer over Unix sockets to bind agent lifecycle events to persistent terminal sessions. By establishing strict bounded I/O budgets, I resolved memory leaks and prevented loops from hanging during high-frequency agent executions.
- **System Observability & Safety (PR #7946):** I built cross-platform CPU/RAM sampling using the sysinfo crate in Rust to profile agent resource consumption, alongside a real-time model-context-window visual status bar in the ZeroCode TUI to proactively prevent context-budget token overflows.
- **QA Automation & AI Evaluation:** Leveraging my background as QA Lead at Knowledgecity, where I managed test automation across 15+ concurrent microservices, I can establish robust, automated verification and evaluation frameworks to benchmark agent behaviors, security boundaries, and data privacy compliance.

Skills Matching & Agentic ELT Architecture Data Flow:



I am highly fluent in leveraging AI in my daily workflows and have built and deployed tools like the herdr-mcp server, which exposes CLI controls to LLMs. I am excited about the prospect of applying my deep systems expertise in Rust, LLM integration, and QA automation to help ClickUp build the next chapter of collaborative multi-agentic workflows.

Sincerely,
Eugene L. Buchanan