

EUGENE L. BUCHANAN

AI Solutions Strategic Consultant — Use Cases to Production ROI

Apple Valley, CA (Remote) · +1 (909) 545 5384 · eugene@serviceofothers.org · github.com/eugeneb50



September 15, 2026

Hiring Team

AI Solutions Strategic Consultant — 100% Remote (12 Months)

Via: requisition portal

Re: AI Solutions Strategic Consultant — 15+ Years, Concept to Production

Dear Hiring Team,

I'm applying for the AI Solutions Strategic Consultant role. For 15+ years I've turned business needs into production systems — most recently into AI solution architectures with clear use cases, technical roadmaps, and measurable ROI. I identify where AI/GenAI pays, evaluate the platforms, models, and vendors honestly, and work with business, data, engineering, and product teams to ship.

For a real estate agency I built and shipped an HNWI finder on county recorder public records. It scores public-records signals into a ranked prospect list with lead enrichment — owner, property, equity, and contactability signals resolved through graph engineering over owners, parcels, and transactions — so agents spend time on the highest-value owners first. The architecture is a data pipeline with API and microservices boundaries feeding existing agency workflows: ingest and normalize public filings, resolve entities, enrich leads, score and rank, and deliver results where agents already work. Retrieval uses RAG patterns with pre- and post-filtering, and failure-model logging feeds a self-improving loop that tightens scoring and enrichment over time. The use case was chosen on value and feasibility — prospecting time is the costliest input — and scoped so it runs on compliant public data within current integrations. That public-records foundation extends my Tierrachain work on parallel records systems and AI-assisted compliance filing for real estate exchange.


I also shipped an eBay listing assistant built on a Hermes agent around a photo upload pipeline: sellers upload item photos, visual reverse image search finds comparable listings, and description generation drafts titles, descriptions, and pricing from the images plus retrieved comparables. The seller problem was concrete: titles, descriptions, and pricing that move inventory without hours of manual drafting. RAG over visual and text comparables applies pre- and post-filtering, the seller stays in the approval loop, and results iterate toward what sells. MCP tool boundaries expose pricing, search, and publishing actions to the agent; failure-model logging captures bad drafts, rejections, and corrections into a self-improving loop. I treated it as a GenAI use case end to end — business value first, then model and platform choice, then evaluation of output quality before release, then hardening from prototype into a tool sellers can use daily.

That product-level delivery rests on agentic infrastructure depth. On ZeroClaw's 32.8k-star Rust gateway I shipped production agent systems: a merged context-window meter unifying 9 LLM providers into one source of truth for model routing and cost attribution (PR #7946), live serving-provider identity on usage events across fallback and model switches (PR #8966, open), and an observability, steering, and interruptible layer over the herdr multiplexer reporting agent lifecycle (idle/working/blocked/released) via JSON-RPC (PR #8337, requirements adopted into the vendor-neutral design) — the coordination plane for A2A-style team frameworks where agents hand off, escalate, and get interrupted safely. At Knowledgecity I owned integrations as Product Owner and shipped SAP, Oracle, Workday, UKG, Coursera, and Zoom over SAML, OAuth, SFTP, and REST APIs, plus a RAG support chatbot with guardrails under SOC 2 controls. I stay current by building where the field moves: MCP servers, graph engineering, and evaluation tooling (Cypress, Selenium, JUnit, dashboards) that proves behavior at the system level.

How I fit this role:

- **AI/GenAI use cases tied to ROI.** HNWI finder with lead enrichment (agent prospecting time) and eBay listing assistant (sell-through per labor hour) — prioritized on value and feasibility.
- **RAG, MCP, graph engineering.** RAG retrieval with pre- and post-filtering; MCP tool boundaries; owner-parcel-transaction graphs for entity resolution and enrichment.
- **Observability, steering, safe interruption.** Herdr-multiplexer lifecycle layer (idle/working/blocked/released) as the coordination plane for A2A-style team handoffs and escalation.
- **Evaluate, learn, ship.** Failure-model logging into a self-improving loop; platform, model, and vendor best-fit analysis; delivery with business, data, engineering, and product teams.

Value: Use Case to Production



I consult the way I build: name the business outcome, map the feasible AI use cases, recommend the simplest architecture that reaches production, and measure whether it paid. I'd welcome a working session on your highest-priority use case — happy to walk through the HNWI scoring pipeline, the listing assistant's evaluation loop, or any ZeroClaw PR.

Best regards,

Eugene L. Buchanan

AI Solutions Strategic Consultant — Concept to Production

Apple Valley, CA (Remote) · eugene@serviceofothers.org · github.com/eugeneb50